

Quantitative Reasoning & Interview Preparation

Informal Notes for Problems on Probability, Estimation, and Financial Intuition

Ryan McMillan, MCompSci (cand.)

Master of Computer Science — University of New England

January 13, 2026

Contents

1	Introduction	2
2	How This Book Is Written	2
3	Coin Flips: Simulation vs Expectation	2

1 Introduction

This document is a curated portfolio of algorithmic solutions. Each selected problem includes:

- a brief problem summary,
- the key idea,
- an algorithm outline,
- a correctness sketch,
- time and space complexity,
- a reference C++ implementation.

The focus is on clarity, correctness, and reasoning.

2 How This Book Is Written

This book is intended to mirror how problems are discussed in quantitative development technical interviews and presented by a candidate on a whiteboard.

The emphasis is on:

- stating assumptions explicitly,
- building intuition before formalism,
- using diagrams where helpful,
- sanity-checking results,
- (optionally) using pseudocode for simulation-style verification.

This is **not** a formal mathematics text. Proofs are informal, notation is kept light, and equations are used only when they improve understanding.

Each entry is a *final explanation*, not a transcript of exploratory work.

3 Coin Flips: Simulation vs Expectation

Problem

Flip a fair coin n times. Let X be the number of heads. What is $\mathbb{E}[X]$? How could we sanity-check the result via simulation?

Assumptions

- The coin is fair: $P(\text{Heads}) = P(\text{Tails}) = \frac{1}{2}$.
- Flips are independent.

Intuition

Each flip contributes either 1 head or 0 heads. On average, a single flip contributes $\frac{1}{2}$ head, so across n flips we expect about $\frac{n}{2}$ heads.

Reasoning

Write $X = \sum_{i=1}^n I_i$, where $I_i = 1$ if flip i is heads and 0 otherwise.

Then:

$$\mathbb{E}[X] = \sum_{i=1}^n \mathbb{E}[I_i]$$

and since $\mathbb{E}[I_i] = P(I_i = 1) = \frac{1}{2}$ for a fair coin:

$$\mathbb{E}[X] = n \cdot \frac{1}{2} = \frac{n}{2}.$$

Procedure (Simulation Pseudocode)

Algorithm 1: Monte Carlo sanity check for $\mathbb{E}[X]$

Input: n flips per trial, T trials

Output: Estimated mean number of heads

sumHeads \leftarrow 0;

for $t \leftarrow 1$ **to** T **do**

heads \leftarrow 0;

for $i \leftarrow 1$ **to** n **do**

 sample $U \sim \text{Uniform}(0, 1)$;

if $U < 1/2$ **then**

heads \leftarrow **heads** + 1;

sumHeads \leftarrow **sumHeads** + **heads**;

return **sumHeads**/ T ;

Sanity Check

- For $n = 1$, $\mathbb{E}[X] = 0.5$ makes sense.
- For large n , the estimate should concentrate near $n/2$.
- If you plot the histogram of heads across trials, it should look roughly binomial.

Takeaway

For n independent fair coin flips, the expected number of heads is:

$$\mathbb{E}[X] = \frac{n}{2}.$$

Simulation provides a quick sanity check without changing the underlying reasoning.